



Experience the latest in app dev at discounted rates.

[See event pricing](#) →



Red Hat OpenShift

Menu ▾

[Home](#) > [Products](#) > [Red Hat OpenShift](#) > [Red Hat OpenShift AI](#)

Red Hat OpenShift AI


[Try it](#)

[Talk to a Red Hatter](#)

What is Red Hat OpenShift AI?

Red Hat® OpenShift® AI is a flexible, scalable MLOps platform with tools to build, deploy, and manage AI-enabled applications. Built using open source technologies, it provides trusted, operationally consistent capabilities for teams to experiment, serve models, and deliver innovative apps.


OpenShift AI (previously called Red Hat OpenShift Data Science) supports the full lifecycle of AI/ML experiments and models, on-premise and in the public cloud.



Launch AI faster in any environment. Video duration: 5:26

Welcome to Red Hat. What brings you here today?

1



Bring AI-enabled apps to production faster

Combine the proven capabilities of Red Hat OpenShift AI and Red Hat OpenShift in a single enterprise-ready AI application platform that brings teams together. Data scientists, engineers, and app developers can collaborate in a single destination that promotes consistency, security, and scalability.

OpenShift AI enables data acquisition and preparation, model training and fine-tuning, model serving and model monitoring, and hardware acceleration. With an open ecosystem of hardware and software partners, OpenShift AI delivers the flexibility you need for your specific use cases.

Read the solution brief →

Use committed AWS spend to purchase and run in AWS Marketplace

Learn more →

Features & benefits



Less time managing AI infrastructure

Provide your teams with on-demand access to resources, so they can focus on exploring data and building apps that add real value to your organization. Additional time-saving benefits include built-in security and operator life cycle integration.

Read the technology overview →



Tested and supported AI/ML tooling

Red Hat tracks, integrates, tests, and supports common AI/ML tooling and model serving on our Red Hat OpenShift application platform, so you don't have to. OpenShift AI draws from years of incubation in Red Hat's Open Data Hub community project.



Flexibility across the hybrid cloud

Offered as either self-managed software or as a fully managed cloud service on top of OpenShift, Red Hat OpenShift AI provides a secure and flexible platform that gives you the choice of where you develop and deploy your models—whether on-premise, the public cloud, or even at the edge.


Discover new features →



Leverage our best practices

Red Hat Consulting and Red Hat Open Innovation Labs provides services that allow you to install, configure and use Red Hat OpenShift AI to its fullest extent.

Learn best practices and build your own ML pipeline project using DevOps and MLOps through Red Hat's proven methodology. Whether you're seeking guidance on building a sample model, or looking to move beyond model experimentation to develop strategies for employing models to production, our comprehensive Services will provide support and mentorship for building modern cloud applications.



Get to know Red Hat OpenShift Service on AWS (ROSA)

Get started

Expand your capabilities with Red Hat's partner ecosystem

Improve cross-functional collaboration and simplify maintenance by extending the core Red Hat OpenShift AI platform with other integrated Red Hat services and products.



Red Hat

OpenShift AI

☰

☐☐☐

🔔

?

demo ▾

Applications ▾

Enabled

Explore

Data Science Projects

Data Science Pipelines >

Model Serving

Resources

Enabled

Launch your enabled applications, view documentation, or get started with quick start instructions and tasks.

jupyter

Jupyter ⚙️

Red Hat managed

A multi-user version of the notebook designed for companies, classrooms and research labs.

Launch application

Collaborating through Jupyter notebooks-as-a-service

Provide pre-built or customized cluster images to your data scientists to build models using Jupyter notebooks. Red Hat OpenShift AI tracks changes to Jupyter, TensorFlow, and PyTorch, and other open source AI technologies.

☰

Red Hat OpenShift AI

Applications

Data Science Projects

Data Science Pipelines

Model Serving

Resources

Settings

Model serving

Manage and view the health and performance of your deployed models.

Name

Find by name

Deploy model

Model name	Project	Model server	Inference endpoint	Status
text_encoder	Text to Image	Custom Model Server	Internal Service	✓
unet	Text to Image	Custom Model Server	Internal Service	✓
vae_decoder	Text to Image	Custom Model Server	Internal Service	✓
vae_encoder	Text to Image	Custom Model Server	Internal Service	✓
yolo	Object Detection	Model Server	Internal Service	✓

Scaling Model Serving with Red Hat OpenShift AI

Models can be served for integration into intelligent applications on-premise, in the public cloud or at the edge. These models can be rebuilt, redeployed, and monitored based on changes to the source notebook.



Innovate with open source and hybrid cloud

Red Hat OpenShift AI integrates open source tools with Red Hat's AI/ML partner ecosystem to accelerate project development and enhance platform capabilities. Plus, with Red Hat's hybrid cloud approach you can eliminate cloud provider lock-ins and minimize vendor lock-ins.

Explore related resources

VIDEO

Demo of Red Hat OpenShift AI

Watch the demo →

WEBINAR

Getting the most out of AI with open source and Kubernetes

Watch the webinar →

CHECKLIST

Top 5 considerations for an AI Platform

Get the checklist →

CHECKLIST

Top 5 ways developers and data scientists can collaborate

Get the checklist →

CONTACT SALES

Talk to a Red Hatter



Products

Tools


Try, buy, & sell

Communicate

About Red Hat

We’re the world’s leading provider of enterprise open source solutions—including Linux, cloud, container, and Kubernetes. We deliver hardened solutions that make it easier for enterprises to work across platforms and environments, from the core datacenter to the network edge.

Select a language

 English ▼



- About Red Hat
- Jobs
- Events
- Locations
- Contact Red Hat
- Red Hat Blog
- Diversity, equity, and inclusion
- Cool Stuff Store
- Red Hat Summit

- © 2024 Red Hat, Inc.
- Privacy statement
- Terms of use
- All policies and guidelines
- Digital accessibility
- Cookie preferences